

- All solutions may be submitted in groups of up to 3 students.
- Send your solutions before the due date per email to `cloudcomputing@uni.lu` in a single zip file (containing just the java files or scripts) using your group id as the file name.
- Include some brief instructions on how to run your solution to each problem in a file called `problem_X.txt`.
- All group members are required to be present in the class for the demonstration of the solutions on the due date.

II. Pig Latin & Hive

For this exercise, please consider downloading the Yago dataset available on the course homepage on `moodle.uni.lu`.

PROCESSING YAGO DATASET

6 Points

Problem 1. Find the top three frequently occurring *predicates* in the Yago dataset using:

- (i) operators available in the Pig Latin scripting language;

3 Points

- (ii) operators available in HiveQL.

3 Points

GROUPING AND JOINING

8 Points

Problem 2. Identify all the given-names (corresponding to `hasGivenName` predicate) of persons who are associated with more than one `livesIn` predicates from the Yago dataset using:

- (i) the relational operations (joins, grouping, etc.) available in the Pig Latin scripting language;

4 Points

- (ii) the relational operations which are available in HiveQL.

4 Points

DATA CLUSTERING & BUCKETIZED MERGE-JOIN

22 Points

Problem 3. Write a HiveQL query to find all the subjects (`x`) and objects (`y` and `z`) matching the pattern: `?x <hasGivenName> ?y. ?x <livesIn> ?z.`, from the Yago dataset.

- Implement this problem by:
 - (i) by considering *partitioning* and *bucketing*;
 - (ii) by considering *partitioning* but not *bucketing*;
 - (iii) by considering neither *partitioning* nor *bucketing*.
- For the first case alone perform a *Bucketized Merge-Join* by enabling the necessary parameters (see the note below).

- Compare the run time of the three cases by performing your experiments on an AWS EC2 instance.
- For case (i):
 - Load the entire triples from the given `yago.tsv` files into a table named `yago` having three columns: *subject*, *predicate* and *object*.
 - Create a new table, named `yago_part_buck`, with a partition based on the predicate column and clustered based on the subject column.
 - Load data (statically) into the partitions for all the 29 predicates¹ in the dataset. This loading could be done by inserting data into the partitioned table from the `yago` table specifying the partition key – you may write all the insert statements into a single HiveQL script for loading data.
 - Write a HiveQL query to find the required pattern from the `yago_part_buck` table.

Note: You can set the following hive parameters to true (as given below), to enable the bucketized merge-join.

```
set hive.auto.convert.sortmerge.join=true;
set hive.optimize.bucketmapjoin = true;
set hive.optimize.bucketmapjoin.sortedmerge = true;
```

¹<actedIn>, <hasAcademicAdvisor>, <hasChild>, <hasFamilyName>, <hasWebsite>, <hasWonPrize>, <isInterestedIn>, <isKnownFor>, <directed>, <edited>, <graduatedFrom>, <hasGender>, <hasMusicalRole>, <isCitizenOf>, <isMarriedTo>, <isPoliticianOf>, <playsFor>, <worksAt>, <wroteMusicFor>, <created>, <diedIn>, <hasGivenName>, <influences>, <isAffiliatedTo>, <isLeaderOf>, <livesIn>, <owns>, <participatedIn>, <wasBornIn>